

Linear regression

Supervised learning example



$a(x)$ → 13

Supervised learning


x_i — example

y_i — target value

$x_i = (x_{i1}, \dots, x_{id})$ — features

$X = ((x_1, y_1), (x_2, y_2), \dots, (x_\ell, y_\ell))$ — training set

$a(x)$ — model, hypothesis



Regression and classification

$y_i \in \mathbb{R}$ — regression task

- Salary prediction
- Movie rating prediction

Regression and classification


$y_i \in \mathbb{R}$ — regression task

- Salary prediction
- Movie rating prediction


y_i belongs to a finite set — classification task

- Object recognition
- Topic classification

Linear model for regression example



Linear model for regression example



Linear model for regression

$$a(x) = b + w_1x_1 + w_2x_2 + \cdots + w_dx_d$$

- w_1, \dots, w_d — coefficients (weights)
- b — bias
- $d + 1$ parameters
- To make it simple: there's always a constant feature

Linear model for regression

Vector notation:

$$a(x) = w^T x$$

For a sample X :

$$a(X) = Xw$$

$$X = \begin{pmatrix} x_{11} & \cdots & x_{1d} \\ \vdots & \ddots & \vdots \\ x_{\ell 1} & \cdots & x_{\ell d} \end{pmatrix}$$

Loss function

How to measure model quality?

Mean squared error:

$$\begin{aligned} L(w) &= \frac{1}{\ell} \sum_{i=1}^{\ell} (w^T x_i - y_i)^2 \\ &= \frac{1}{\ell} \|Xw - y\|^2 \end{aligned}$$

Training a model

Fitting a model to training data:

$$L(w) = \frac{1}{\ell} \|Xw - y\|^2 \rightarrow \min_w$$

Training a model

Fitting a model to training data:

$$L(w) = \frac{1}{\ell} \|Xw - y\|^2 \rightarrow \min_{w,}$$

Exact solution:

$$w = (X^T X)^{-1} X^T y$$

But inverting a matrix is hard for high-dimensional data!

Summary

- Linear models are very simple
- MSE can be used as a loss function
- There is an analytical solution, but we need more generic and scalable learning method